

Import Packages

```
1 import numpy as np
2 import pandas as pd
3 import matplotlib.pyplot as plt
4 import seaborn as sns
5
6 # Data preprocessing
7 from sklearn.preprocessing import MinMaxScaler
8 from sklearn.model_selection import train_test_split
9
10 # Neural net model
11 from keras.models import Sequential
12 from keras.layers import Dense, Dropout
13 from keras.callbacks import EarlyStopping
```

Read in the data

```
1 df = pd.read_csv('housing.csv')
2 df.head()
```

	longitude	latitude	housing_median_age	total_rooms	total_bedrooms	population	households	median_income	median_house_value
0	-122.23	37.88	41.0	880.0	129.0	322.0	126.0	8.3252	452600.0
1	-122.22	37.86	21.0	7099.0	1106.0	2401.0	1138.0	8.3014	358500.0
2	-122.24	37.85	52.0	1467.0	190.0	496.0	177.0	7.2574	352100.0
3	-122.25	37.85	52.0	1274.0	235.0	558.0	219.0	5.6431	341300.0
4	-122.25	37.85	52.0	1627.0	280.0	565.0	259.0	3.8462	342200.0

Data Preprocessing and cleaning

```
1 df.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 20640 entries, 0 to 20639
Data columns (total 10 columns):
#   Column                Non-Null Count  Dtype
---  -
0   longitude              20640 non-null float64
1   latitude               20640 non-null float64
2   housing_median_age     20640 non-null float64
3   total_rooms            20640 non-null float64
4   total_bedrooms         20433 non-null float64
5   population             20640 non-null float64
6   households              20640 non-null float64
7   median_income          20640 non-null float64
8   median_house_value     20640 non-null float64
9   ocean_proximity        20640 non-null object
dtypes: float64(9), object(1)
memory usage: 1.6+ MB
```

```
1 df.isnull().sum()
```

```
longitude      0
latitude       0
housing_median_age  0
total_rooms    0
total_bedrooms 207
population     0
households     0
median_income  0
median_house_value  0
ocean_proximity  0
dtype: int64
```

```
1 # drop rows with missing values
2 df.dropna(axis=0, inplace=True)
```

```
1 df.shape
```

(20433, 10)

```
1 df['ocean_proximity'].value_counts()
```

```
ocean_proximity
<1H OCEAN    9034
INLAND       6496
NEAR OCEAN   2628
NEAR BAY     2270
ISLAND        5
Name: count, dtype: int64
```

```
1 df['ocean_proximity'] = df['ocean_proximity'].map({'<1H OCEAN':0,"INLAND":1,"NEAR OCEAN":2,"NEAR BAY":3,"ISLAND":4})
```

```
1 df.head()
```

```
longitude latitude housing_median_age total_rooms total_bedrooms population
0 -122.23 37.88 41.0 880.0 129.0 322.0
1 -122.22 37.86 21.0 7099.0 1106.0 2401.0
2 -122.24 37.85 52.0 1467.0 190.0 496.0
3 -122.25 37.85 52.0 1274.0 235.0 558.0
4 -122.25 37.85 52.0 1627.0 280.0 565.0
```

✓ Train test split

```
1 # Target variable is 'median_house_value'
2 y = df['median_house_value']
3 x = df.drop('median_house_value', axis=1)
4 print(x.shape)
5 print(y.shape)
```

(20433, 9)
(20433,)

```
1 # convert to numpy array
2 x = np.array(x)
3 y = np.array(y)
```

```
1 x_train, x_test, y_train, y_test = train_test_split(x, y, test_size=0.2, random_state=123)
2
3 print(x_train.shape)
4 print(x_test.shape)
5 print(y_train.shape)
6 print(y_test.shape)
```

(16346, 9)
(4087, 9)
(16346,)
(4087,)

✓ Scaling

```
1 # Use MinMax scaler
2 min_max_scaler = MinMaxScaler()
3 x_train = min_max_scaler.fit_transform(x_train)
4 x_test = min_max_scaler.transform(x_test)
```

```
1 print(x_train)
2 print(x_test)
```

```
[[0.54425229 0.30180659 0.66666667 ... 0.06101974 0.21336947 0.25 ]
 [0.30315361 0.58448459 0.56862745 ... 0.00707237 0.27586516 0.25 ]
 [0.6948118 0.16790648 0.09803922 ... 0.25542763 0.25377581 0.25 ]
 ...
 [0.23601221 0.51328374 0.68627451 ... 0.05411184 0.2810375 0. ]
 [0.19226857 0.55366631 1. ... 0.06069079 0.26132053 0.75 ]
 [0.51068159 0.19978746 0.43137255 ... 0.12648026 0.29727866 0.5 ]
 [[0.71922686 0.02231668 0.43137255 ... 0.09769737 0.18199749 0.5 ]
 [0.30620549 0.64718385 0.45098039 ... 0.07976974 0.19436284 0.25 ]]
```

```
[0.3387589 0.54091392 0.74509804 ... 0.09917763 0.08313678 0.25      ]
...
[0.15055951 0.62274176 0.76470588 ... 0.06973684 0.23342437 0.      ]
[0.6103764 0.17534538 1.          ... 0.04276316 0.29645798 0.      ]
[0.72227874 0.08926674 0.25490196 ... 0.13289474 0.31384395 0.      ]]
```


```
1 x_train.shape[1]
```

 9

✓

Building model (ANN)


```
1 model = Sequential([
2     #input layer
3     Dense(1000, input_shape=(x_train.shape[1],), activation='relu'),
4     Dropout(0.2),
5     #two hidden layers
6     Dense(500, activation='relu'),
7     Dropout(0.2),
8     Dense(250, activation='relu'),
9     #output layer
10    Dense(1,activation='linear') # here 1 shows countinuous value(regression)
11 ])
12 model.summary()
```

 Model: "sequential"

Layer (type)	Output Shape	Param #
=====		
dense (Dense)	(None, 1000)	10000
dropout (Dropout)	(None, 1000)	0
dense_1 (Dense)	(None, 500)	500500
dropout_1 (Dropout)	(None, 500)	0
dense_2 (Dense)	(None, 250)	125250
dense_3 (Dense)	(None, 1)	251
=====		
Total params: 636001 (2.43 MB)		
Trainable params: 636001 (2.43 MB)		
Non-trainable params: 0 (0.00 Byte)		

```
1 # Compile the model and set easlystopping
2 model.compile(optimizer='rmsprop', loss='mse', metrics=['mae'])
3 es = EarlyStopping(monitor='val_loss',mode='min',patience=50, restore_best_weights=True)
```

```
1 # fit the model (training)
2 history = model.fit(x_train, y_train, validation_data= (x_test, y_test), callbacks=[es], epochs=10, batch_size=50, verbose=1)
```

 Epoch 1/10
327/327 [=====] - 7s 19ms/step - loss: 24755955712.0000 - mae: 120739.8438 - val_loss: 10583593984.0000 - val_

Epoch 2/10
327/327 [=====] - 5s 15ms/step - loss: 9480070144.0000 - mae: 77307.1484 - val_loss: 7265815040.0000 - val_

Epoch 3/10
327/327 [=====] - 6s 20ms/step - loss: 6495425024.0000 - mae: 60929.1836 - val_loss: 5670961664.0000 - val_

Epoch 4/10
327/327 [=====] - 5s 16ms/step - loss: 5929728000.0000 - mae: 56496.6914 - val_loss: 5567024640.0000 - val_

Epoch 5/10
327/327 [=====] - 5s 14ms/step - loss: 5745964544.0000 - mae: 55423.3672 - val_loss: 5320379904.0000 - val_

Epoch 6/10
327/327 [=====] - 6s 19ms/step - loss: 5555000832.0000 - mae: 54292.5625 - val_loss: 5109794816.0000 - val_

Epoch 7/10
327/327 [=====] - 5s 14ms/step - loss: 5373375488.0000 - mae: 53132.9141 - val_loss: 4974953984.0000 - val_

Epoch 8/10
327/327 [=====] - 6s 17ms/step - loss: 5202966528.0000 - mae: 52000.4180 - val_loss: 4817773056.0000 - val_

Epoch 9/10
327/327 [=====] - 10s 30ms/step - loss: 5059791872.0000 - mae: 51048.3125 - val_loss: 4650961920.0000 - val_

Epoch 10/10
327/327 [=====] - 8s 23ms/step - loss: 4887821312.0000 - mae: 50195.3477 - val_loss: 4527864832.0000 - val_

✓

Metrics and Score

```

1 from sklearn.metrics import r2_score, mean_absolute_error, mean_squared_error, mean_squared_log_error
2
3 y_pred = model.predict(x_test)
4 print("mae :", mean_absolute_error(y_test, y_pred))
5 print("mse :", mean_squared_error(y_test, y_pred))
6 print("mae :", mean_squared_log_error(y_test, y_pred))
7 print("score :", r2_score(y_test, y_pred))

```

```


128/128 [=====] - 1s 5ms/step
mae : 47538.12862334077
mse : 4527865104.3238945
mae : 0.10731578783218153
score : 0.6468630546819467

```

```

1 # Get the training and validation loss from the history object
2 training_loss = history.history['loss']
3 validation_loss = history.history['val_loss']
4 epochs = range(1, len(training_loss) + 1)
5 # Plot the training and validation loss
6 plt.plot(epochs, training_loss, 'b', label='Training Loss')
7 plt.plot(epochs, validation_loss, 'r', label='Validation Loss')
8
9 # Label the plot
10 plt.title('Training and Validation Loss')
11 plt.xlabel('Epochs')
12 plt.ylabel('Loss')
13 plt.legend()
14
15 # Show the plot
16 plt.show()

```



```

1 # Specify the dimensions
2 fig, axes = plt.subplots(1,2)
3
4 # This makes the individual subplots
5 # Training Results
6 axes[0].scatter(x=y_train, y=model.predict(x_train))
7 axes[0].set_xlabel('Actual', fontsize=10)
8 axes[0].set_ylabel('Prediction', fontsize=10)
9 axes[0].set_title('Training')
10
11 # Add 45 deg line
12 x = np.linspace(*axes[0].get_xlim())
13 axes[0].plot(x, x, color='red')
14
15 # Validation Results
16 axes[1].scatter(x=y_test, y=model.predict(x_test))
17 axes[1].set_xlabel('Actual', fontsize=10)
18 axes[1].set_ylabel('Predicted', fontsize=10)
19 axes[1].set_title('Validation')
20
21 # add 45 deg line
22 x = np.linspace(*axes[1].get_xlim())
23 axes[1].plot(x, x, color='red')
24
25 # Tight layout
26 fig.tight_layout()
27 plt.show()

```

511/511 [=====] - 2s 4ms/step
128/128 [=====] - 1s 7ms/step

✓ Predictive System

```

1 def pred(longitude,latitude,housing_median_age,total_rooms,total_bedrooms,population,households,median_income,ocean_proximity):
2     features = np.array([longitude,latitude,housing_median_age,total_rooms,total_bedrooms,population,households,median_income,ocean_
3     features_scaled = min_max_scaler.fit_transform([features])
4     results = model.predict(features_scaled).reshape(1,-1)
5     return results[0]

```

longitude = -122.2300
latitude = 37.8800
housing_median_age = 41.0000
total_rooms = 880.0000
total_bedrooms = 129.0000
population = 322.0000
households = 126.0000
median_income = 8.3252
ocean_proximity = 3.0000

```

1 price = pred(longitude,latitude,housing_median_age,total_rooms,total_bedrooms,population,households,median_income,ocean_proximity)
12

```

1/1 [=====] - 0s 31ms/step

1 price #452600.0000 #actual value

array([270706.47], dtype=float32)